

Quora Question Pairs

1. Business Problem

1.1 Description

Quora is a place to gain and share knowledge—about anything. It's a platform to ask questions and connect with people who contribute unique insights and quality answers. This empowers people to learn from each other and to better understand the world.

Over 100 million people visit Quora every month, so it's no surprise that many people ask similarly worded questions. Multiple questions with the same intent can cause seekers to spend more time finding the best answer to their question, and make writers feel they need to answer multiple versions of the same question. Quora values canonical questions because they provide a better experience to active seekers and writers, and offer more value to both of these groups in the long term.

Credits: Kaggle

Problem Statement

- Identify which questions asked on Quora are duplicates of questions that have already been asked.
- This could be useful to instantly provide answers to questions that have already been answered.
- We are tasked with predicting whether a pair of questions are duplicates or not.

1.3 Real world/Business Objectives and Constraints

1. The cost of a mis-classification can be very high.
2. You would want a probability of a pair of questions to be duplicates so that you can choose any threshold of choice.
3. No strict latency concerns.
4. Interpretability is partially important.

2. Machine Learning Problem

2.1 Data

2.1.1 Data Overview

- Data will be in a file Train.csv
- Train.csv contains 5 columns : qid1, qid2, question1, question2, is_duplicate
- Size of Train.csv - 60MB
- Number of rows in Train.csv = 404,290

2.2 Mapping the real world problem to an ML problem

2.2.2 Performance Metric

Source: <https://www.kaggle.com/c/quora-question-pairs#evaluation>

Metric(s):

- log-loss : <https://www.kaggle.com/wiki/LogarithmicLoss>
- Binary Confusion Matrix

1. Business Problem

```
In [1]: from google.colab import drive
drive.mount('/content/drive')
```

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly

Enter your authorization code:

.....

Mounted at /content/drive

```
In [2]: pip install git+git://github.com/seatgeek/fuzzywuzzy.git@0.17.0#egg=fuzzywuzzy
```

Collecting fuzzywuzzy

Cloning git://github.com/seatgeek/fuzzywuzzy.git (to revision 0.17.0) to /tmp/pip-install-gdej9nyx/fuzzywuzzy

Running command git clone -q git://github.com/seatgeek/fuzzywuzzy.git /tmp/pip-install-gdej9nyx/fuzzywuzzy

Running command git checkout -q 089e5c46a117e3de293fd6cb94580efc2b9f6912

Building wheels for collected packages: fuzzywuzzy

Building wheel for fuzzywuzzy (setup.py) ... done

Created wheel for fuzzywuzzy: filename=fuzzywuzzy-0.17.0-py2.py3-none-any.whl size=18151 sha256=85b2fe8b0959d0ae85d04e142bdbdfd8c1c3246924399f413826a2081e55d1b6

Stored in directory: /tmp/pip-ephem-wheel-cache-3zul206g/wheels/93/cb/79/4e772f8c8772e2f67a1548610943f8b23b2a83a278206bc5d9

Successfully built fuzzywuzzy

Installing collected packages: fuzzywuzzy

Successfully installed fuzzywuzzy-0.17.0

```
In [3]: '!python -m spacy download en_core_web_lg'''
```

```
Out[3]: '!python -m spacy download en_core_web_lg'
```

```
In [4]: pip install distance
```

```
Collecting distance
  Downloading https://files.pythonhosted.org/packages/5c/1a/883e47df323437aef
fa0d0a92ccfb38895d9416bd0b56262c2e46a47767b8/Distance-0.1.3.tar.gz (180kB)
    |██████████████████████████████████████| 184kB 2.8MB/s
Building wheels for collected packages: distance
  Building wheel for distance (setup.py) ... done
  Created wheel for distance: filename=Distance-0.1.3-cp36-none-any.whl size
=16261 sha256=92ad110a8e5347b74b531f890094a18119235757d8c6387c16a4fc3e07e9da
e5
  Stored in directory: /root/.cache/pip/wheels/d5/aa/e1/dbba9e7b6d397d645d0f
12db1c66dbae9c5442b39b001db18e
Successfully built distance
Installing collected packages: distance
Successfully installed distance-0.1.3
```

```
In [5]: pip install python-Levenshtein
```

```
Collecting python-Levenshtein
  Downloading https://files.pythonhosted.org/packages/42/a9/d1785c85ebf9b7dfacd08938dd028209c34a0ea3b1bcdb895208bd40a67d/python-Levenshtein-0.12.0.tar.gz (48kB)
    |████████████████████| 51kB 1.6MB/s
Requirement already satisfied: setuptools in /usr/local/lib/python3.6/dist-packages (from python-Levenshtein) (41.4.0)
Building wheels for collected packages: python-Levenshtein
  Building wheel for python-Levenshtein (setup.py) ... done
  Created wheel for python-Levenshtein: filename=python_Levenshtein-0.12.0-cp36-cp36m-linux_x86_64.whl size=144669 sha256=44b93a21b0c30ac6a7d4371dcd3fc3e6ed6c3cfa94704e097f3a98b41514d2a2
  Stored in directory: /root/.cache/pip/wheels/de/c2/93/660fd5f7559049268ad2dc6d81c4e39e9e36518766eaf7e342
Successfully built python-Levenshtein
Installing collected packages: python-Levenshtein
Successfully installed python-Levenshtein-0.12.0
```

```
In [6]: %matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import string
import nltk
nltk.download('stem')
from nltk.stem import PorterStemmer
from bs4 import BeautifulSoup
from fuzzywuzzy import fuzz
import distance
from fuzzywuzzy import process
import os
from scipy.sparse import hstack as hs

# for text featurization
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
import spacy
from sklearn.model_selection import train_test_split

from numpy import hstack

# model
from sklearn.linear_model import SGDClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics.classification import accuracy_score, log_loss
```

```

from collections import Counter
from sklearn.svm import SVC
from sklearn.model_selection import StratifiedKFold
from collections import Counter, defaultdict
from sklearn.calibration import CalibratedClassifierCV
from sklearn.naive_bayes import MultinomialNB
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV, RandomizedSearchCV
import math
from sklearn.metrics import normalized_mutual_info_score

from sklearn.model_selection import cross_val_score
from mlxtend.classifier import StackingClassifier
from xgboost import XGBClassifier

from sklearn import model_selection
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import precision_recall_curve, auc, roc_curve

import plotly.offline as offline
import plotly.graph_objs as go
offline.init_notebook_mode()
import numpy as np
import matplotlib.patches as mpatches

import warnings
warnings.filterwarnings("ignore")

```

```
[nltk_data] Error loading stem: Package 'stem' not found in index
```

```

/usr/local/lib/python3.6/dist-packages/sklearn/externals/six.py:31: DeprecationWarning: The module is deprecated in version 0.21 and will be removed in version 0.23 since we've dropped support for Python 2.7. Please rely on the official version of six (https://pypi.org/project/six/).
  "(https://pypi.org/project/six/).", DeprecationWarning)

```

```

In [0]: data = pd.read_csv('/content/drive/My Drive/Quora/train.csv')
data = data.sample(100000)

```

3. Exploratory Data Analysis

```
In [161]: data.columns
```

```
Out[161]: Index(['id', 'qid1', 'qid2', 'question1', 'question2', 'is_duplicate'], dtype='object')
```

```
In [162]: data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 100000 entries, 254740 to 232193
Data columns (total 6 columns):
id                100000 non-null int64
qid1              100000 non-null int64
qid2              100000 non-null int64
question1         100000 non-null object
question2         100000 non-null object
is_duplicate      100000 non-null int64
dtypes: int64(4), object(2)

```

```
In [163]: data.head()
```

```
Out[163]:
```

	id	qid1	qid2	question1	question2	is_duplicate
254740	254740	369525	369526	How can I stop myself thinking about my ex who...	How do I stop myself from contacting/thinking ...	1
341718	341718	418224	294841	How can I make ice cream project?	How do you make ice cream?	0
293090	293090	48667	278028	What is best website for learning?	What is the best website for online learning?	1
165871	165871	83026	86107	Why did MS Dhoni give up captaincy?	Why MS Dhoni has quit the captaincy from limit...	1
183808	183808	280976	280977	What builds work well in Protoss vs Zerg in Le...	StarCraft II: What are some of the easiest way...	0

Distribution of data points among output classes

- Number of duplicate(smilar) and non-duplicate(non similar) questions

```
In [164]: total = data.shape[0]
duplicate = data[data.is_duplicate==1].shape[0]
not_duplicate = data[data.is_duplicate==0].shape[0]

print("similar question pairs {:.2f}%".format((duplicate/total)*100))
print("non similar question pairs {:.2f}%".format((not_duplicate/total)*100))

similar question pairs 37.09%
non similar question pairs 62.91%
```

Above we can see that only 36.92% of question are duplicate and 63.08% percent are not duplicate. that means there are more number of non duplicate questions.

```
In [165]: data.groupby('is_duplicate')['id'].count().plot.bar()
plt.title('bar graph representing duplicate and non duplicate questions')
```

```
Out[165]: Text(0.5, 1.0, 'bar graph representing duplicate and non duplicate question
s')
```

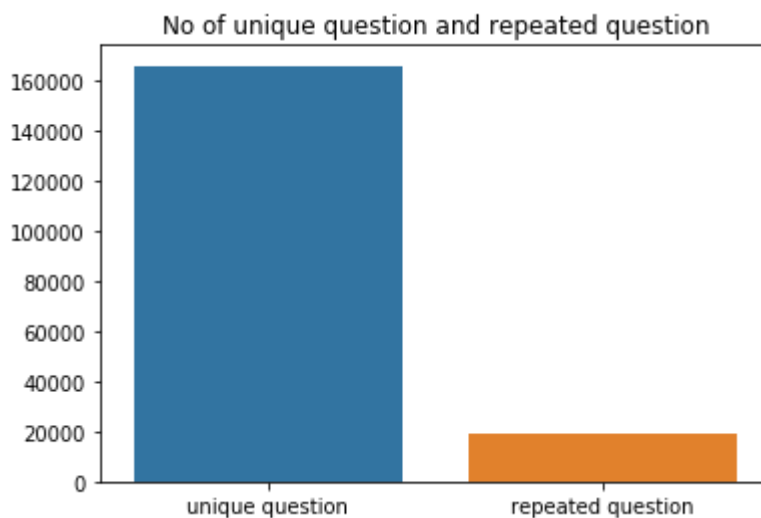
bar graph representing duplicate and non duplicate questions

Number of unique questions

```
In [166]: questions = pd.Series(data['qid1'].tolist() + data['qid2'].tolist())
unique = len(np.unique(questions))
more_than_1 = np.sum(questions.value_counts()>1)
print("total number of unique questions are = ", unique)
print("total number of questions appearing more than once = ", more_than_1)

total number of unique questions are = 166153
total number of questions appearing more than once = 19316
```

```
In [167]: sns.barplot(['unique question', 'repeated question'], [unique, more_than_1])
plt.title('No of unique question and repeated question')
plt.show()
```



checking for null values

```
In [168]: # now looking for null values in the data
print("is there null value in question1= {} , {} questions are null".format(data['question1'].isnull().any(), data['question1'].isnull().sum() ))
print("is there null value in question2= {}, {} questions are null".format(data['question2'].isnull().any(), data['question2'].isnull().sum()))

is there null value in question1= False , 0 questions are null
is there null value in question2= False, 0 questions are null
```

```
In [0]: # we can see that both question1 and question2 can have null values
# filling null values with ''
data['question1'].fillna(' ', inplace=True)
data['question2'].fillna(' ', inplace=True)
```

```
In [170]: print("is there null value in question1= {} , {} questions are null".format(data['question1'].isnull().any(), data['question1'].isnull().sum() ))
print("is there null value in question2= {}, {} questions are null".format(data['question2'].isnull().any(), data['question2'].isnull().sum()))

is there null value in question1= False , 0 questions are null
is there null value in question2= False, 0 questions are null
```


Number of occurrences of each question

```
In [0]: # getting the number of occurrence of each question
qids = []
qids.extend(data['qid1'].tolist())
qids.extend(data['qid2'].tolist())

from collections import Counter

occurrence = Counter(qids)
```

```
In [172]: %%time
plt.figure(figsize=(10,5))
plt.bar(list(occurrence.values()),range(len(occurrence)), align='center')
plt.xlabel('number of occurrence')
plt.ylabel('number of questions')
plt.title('number of occurrence of each question')
plt.show()
plt.close()
```



CPU times: user 5min 53s, sys: 2.27 s, total: 5min 55s
Wall time: 5min 55s

```
In [0]: # getting frequency of each questions
q1 = Counter(data['qid1'].tolist())
q2 = Counter(data['qid2'].tolist())

data['freq_1'] = [q1[i] for i in data['qid1'].tolist() ]
data['freq_2'] = [q2[i] for i in data['qid2'].tolist() ]
```

```
In [174]: # removing stopwords from the question
# https://www.geeksforgeeks.org/removing-stop-words-nltk-python/
import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
from tqdm import tqdm

stop_words = set(stopwords.words('english'))

filtered_isduplicate1 = []
filtered_isduplicate2 = []
filtered_notduplicate1 = []
filtered_notduplicate2 = []
```

```

for i, j in tqdm(zip(data[data.is_duplicate==0]['question1'].tolist(), data[dat
a.is_duplicate==0]['question2'].tolist())):
    temp = ''
    for w in i.split():
        if w not in stop_words:
            temp = temp + w + " "
    filtered_notduplicate1.append(temp)

    temp = ''
    for w in j.split():
        if w not in stop_words:
            temp = temp + w + " "
    filtered_notduplicate2.append(temp)

for i, j in tqdm(zip(data[data.is_duplicate==1]['question1'].tolist(), data[dat
a.is_duplicate==1]['question2'].tolist())):
    temp = ''
    for w in i.split():
        if w not in stop_words:
            temp = temp + w + " "
    filtered_isduplicate1.append(temp)

    temp = ''
    for w in j.split():
        if w not in stop_words:
            temp = temp + w + " "
    filtered_isduplicate2.append(temp)

```

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!

62910it [00:00, 120167.49it/s]
37090it [00:00, 138099.01it/s]

Word cloud for common words

In [0]: *# getting common words list from the questions*

```

common_words_isduplicate = ''
common_words_notduplicate = ''
for i, j in zip(filtered_isduplicate1, filtered_isduplicate2):
    for w in i.split():
        if w in j:
            common_words_isduplicate += w + " "

for i, j in zip(filtered_notduplicate1, filtered_notduplicate2):
    for w in i.split():
        if w in j:
            common_words_notduplicate += w + " "

```

In [176]: *# generating word cloud for similar questions*

```

# https://www.geeksforgeeks.org/generating-word-cloud-python/
from wordcloud import WordCloud, STOPWORDS

stopwords = set(STOPWORDS)

wordcloud = WordCloud(width = 800, height = 800,

```




```

        background_color = 'white',
        stopwords = stopwords,
        min_font_size = 10).generate(common_words_isduplicate)

# plot the WordCloud image
plt.figure(figsize = (8, 8), facecolor = None)
plt.imshow(wordcloud)
plt.axis("off")
plt.title("WordCloud for common word for duplicate question")
plt.tight_layout(pad = 0)

plt.show()

```



Word cloud for dissimilar texts

```
In [177]: # generating word cloud for dissimilar question
wordcloud = WordCloud(width = 800, height = 800,
                      background_color = 'white',
                      stopwords = stopwords,
```


				cream project?				
293090	293090	48667	278028	What is best website for learning?	What is the best website for online learning?	1	1	2
165871	165871	83026	86107	Why did MS Dhoni give up captaincy?	Why MS Dhoni has quit the captaincy from limit...	1	2	5
183808	183808	280976	280977	What builds work well in Protoss vs Zerg in Le...	StarCraft II: What are some of the easiest way...	0	1	1

Preprocessing the data

```
In [0]: def preprocess(text):
'''
function to preprocess each question
It includes:-
1. Removing html tags
2. removing punctuations
3. removing stemming
4. removing stopwords
'''
# removing html tags using beautifulsoup4
text = text.lower()
text = BeautifulSoup(text)
text = text.get_text()

# removing punctuations for the questions
# https://stackoverflow.com/a/16799238/11746488

text = re.sub(r'^\w\s0-9',' ',text)

# removing stemming words and stop words

pattern = re.compile('\W')
text = re.sub(pattern, ' ', text)

stemming = PorterStemmer()
temp_text = ''
for word in text.split():
    if (word not in stop_words) and (type(word) == type('')):
        temp_text += stemming.stem(word) + ' '
return temp_text.strip(' ')
```

```
In [180]: # getting preprocessed text in a list
question1 = []
question2 = []
for i in tqdm(data['question1'].tolist()):
    question1.append(preprocess(i))
for i in tqdm(data['question2'].tolist()):
    question2.append(preprocess(i))
```

```
100%|██████████| 100000/100000 [00:45<00:00, 2210.72it/s]
100%|██████████| 100000/100000 [00:45<00:00, 2190.72it/s]
```

```
In [0]: def extract_features(text1, text2, i, a):
        '''
        funtion to extract additional feature
        it includes:- 1. token count, 2. counting common token ratio, 3. counting s
topword ratio
                        4. common word count, 5. avg token length, 6. absolute length
difference,
                        7. common first words, 8. fuzzy ratio, 9. fuzzy token set,
                        10. fuzzy token sort, 11. fuzzy partial ratio, 12, longest co
mmon sequence

        '''

        token1 = text1.lower().split()
        token2 = text2.lower().split()

        # counting common token ratio
        # https://www.geeksforgeeks.org/python-intersection-of-two-string/
        # using set() + intersection() to
        # get string intersection
        ctc = len(set(token1).intersection(token2))
        if token1>token2:
            try:
                a[0][i] = (float(ctc/len(token2)))
            except ZeroDivisionError:
                a[0][i] = 0

            try:
                a[1][i] = (float(ctc/len(token1)))
            except ZeroDivisionError:
                a[1][i] = 0

        else:
            try:
                a[0][i] = (float(ctc/len(token1)))
            except ZeroDivisionError:
                a[0][i] = 0

            try:
                a[1][i] = (float(ctc/len(token2)))
            except ZeroDivisionError:
                a[1][i] = 0

        # common stop count
        q1 = [w for w in text1.split() if w in stop_words]
        q2 = [w for w in text2.split() if w in stop_words]

        csw = len(set(q1).intersection(q2))
        if len(q1)>len(q2):
            try:
                a[2][i] = (float(csw/len(q2)))
            except ZeroDivisionError:
                a[2][i] = 0

            try:
                a[3][i] = (float(csw/len(q1)))
            except ZeroDivisionError:
                a[3][i] = 0
        else:
            try:
                a[2][i] = (float(csw/len(q1)))
            except ZeroDivisionError:
                a[2][i]=0
```

```

        try:
            a[3][i] = (float(csw/len(q2)))
        except ZeroDivisionError:
            a[3][i] = 0

# common word count
q1 = [w for w in text1.split() if w not in stop_words]
q2 = [w for w in text2.split() if w not in stop_words]

cwc = len(set(q1).intersection(q2))
if len(q1)>len(q2):
    try:
        a[4][i] = (float(cwc/len(q2)))
    except ZeroDivisionError:
        a[4][i] = 0
    try:
        a[5][i] = (float(cwc/len(q1)))
    except ZeroDivisionError:
        a[5][i] = 0
else:
    try:
        a[4][i] = (float(cwc/len(q1)))
    except ZeroDivisionError:
        a[4][i]=0
    try:
        a[5][i] = (float(cwc/len(q2)))
    except ZeroDivisionError:
        a[5][i]=0

# comparing common first word
try:
    a[6][i] = (int( text1.split()[0]== text2.split()[0]))
except:
    a[6][i] = 0

# average token length
a[7][i] = ((len(token1)+len(token2))/2)

# absolute length difference b/w two questions
a[8][i] = (abs(len(token1) - len(token2)))

# levenstein distance
''' In information theory, linguistics and computer science, the Levenshtein distance is a string metric for measuring the difference between two sequences. Informally, the Levenshtein distance between two words is the minimum number of single-character edits (insertions, deletions or substitutions) required to change one word into the other.'''
# fuzzy ratio
a[9][i] = (fuzz.ratio(text1.lower(), text2.lower()))

# fuzzy partial ratio
a[10][i] = (fuzz.partial_ratio(text1.lower(), text2.lower()))

# fuzzy token sort ratio
a[11][i] = (fuzz.token_sort_ratio(text1.lower(), text2.lower()))

# fuzzy token set ratio
a[12][i] = (fuzz.token_set_ratio(text1.lower(), text2.lower()))

# longest substring ratio
lcs = list(distance.lcs substrings(text1, text2))
if len(lcs) == 0:
    a[13][i] = (0)
else:

```

```
a[13][i] = (len(lcs) / (min(len(token1), len(token2)) + 1))
```

```
return a
```

```
In [0]: def add_features():

    a = np.zeros((14, len(data)))

    text1 = data['question1'].tolist()
    text2 = data['question2'].tolist()

    for i in tqdm(range(len(data))):
        a = extract_features(text1[i], text2[i], i, a)
        i+=1

    data['ctc_min'] = a[0]
    data['ctc_max'] = a[1]
    data['csw_min'] = a[2]
    data['csw_max'] = a[3]
    data['cwc_min'] = a[4]
    data['cwc_max'] = a[5]
    data['first_word'] = a[6]
    data['avg_token_length'] = a[7]
    data['abs_lnth_diff'] = a[8]
    data['fzy_ratio'] = a[9]
    data['fzy_partial_ratio'] = a[10]
    data['fzy_token_set'] = a[11]
    data['fzy_token_sort'] = a[12]
    data['lcsr'] = a[13]
```

```
In [183]: add_features()

100%|██████████| 100000/100000 [01:31<00:00, 1095.31it/s]
```

```
In [0]: # adding preprocessed text to data set
data['preprocessed_q1'] = question1
data['preprocessed_q2'] = question2
```

```
In [185]: data.head()
```

Out[185]:

	id	qid1	qid2	question1	question2	is_duplicate	freq_1	freq_
254740	254740	369525	369526	How can I stop myself thinking about my ex who...	How do I stop myself from contacting/thinking ...	1	1	1
341718	341718	418224	294841	How can I make ice cream project?	How do you make ice cream?	0	1	1
293090	293090	48667	278028	What is best website for learning?	What is the best website for online learning?	1	1	2
165871	165871	83026	86107	Why did MS Dhoni	Why MS Dhoni has quit the	1	2	5

				give up captaincy?	captaincy from limit...			
183808	183808	280976	280977	What builds work well in Protoss vs Zerg in Le...	StarCraft II: What are some of the easiest way...	0	1	1

```
In [0]: y = data['is_duplicate']
x = data.drop(['is_duplicate', 'question1', 'question2'], axis=1)
Xtrain, Xtest, Ytrain, Ytest = train_test_split(x, y, test_size=0.33)
```

```
In [248]: count_ = list(Ytrain).count(0)
print("distribution of non duplicate class ", count_/len(Ytrain))
print("distribution of duplicate class ", (len(Ytrain)-count_)/len(Ytrain))

distribution of non duplicate class  0.6291492537313433
distribution of duplicate class  0.3708507462686567
```

```
In [0]: # featurizing text data with tfidf-word vector
questions = list(Xtrain['preprocessed_q1']) + list(Xtrain['preprocessed_q1'])
tfidf = TfidfVectorizer()
tfidf.fit(questions)
a = (tfidf.transform(Xtrain['preprocessed_q1']))
b = (tfidf.transform(Xtrain['preprocessed_q2']))

c = tfidf.transform(Xtest['preprocessed_q1'])
d = tfidf.transform(Xtest['preprocessed_q2'])

word2tfidf = dict(zip(tfidf.get_feature_names(), tfidf.idf_))
```

```
In [0]: def w2vec(questions, word2tfidf):
    nlp = spacy.load('en_core_web_sm')
    vector = []
    for q in tqdm(questions):
        doc = nlp(q) # getting word representation of the question
        try:
            mean_vector = np.zeros((len(doc), len(doc[0].vector)))
        except:
            doc = nlp(' ')
            mean_vector = np.zeros((len(doc), len(doc[0].vector)))
            vector.append(mean_vector.mean(axis=0))
            continue
        for word in doc:
            vect = word.vector # getting the vector representation for each word

            # getting idf values of each word from tfidf vector
            try:
                idf = word2tfidf[str(word)]
            except:
                idf = 0
            mean_vector += vect * idf # multiplying vector representation of the word with its idf values
        mean_vector = mean_vector.mean(axis=0)
        vector.append(mean_vector)
    return vector
```

```
In [251]: # getting w2vec vector for train data
```

```

a1= w2vec(list(Xtrain['preprocessed_q1']), word2tfidf)

a2= w2vec(list(Xtrain['preprocessed_q2']), word2tfidf)

100%|██████████| 67000/67000 [09:53<00:00, 112.81it/s]
100%|██████████| 67000/67000 [09:55<00:00, 112.56it/s]

```

```

In [252]: # getting word to vec for test data
b1 = w2vec(list(Xtest['preprocessed_q1']), word2tfidf)

b2 = w2vec(list(Xtest['preprocessed_q2']), word2tfidf)

100%|██████████| 33000/33000 [04:54<00:00, 112.06it/s]
100%|██████████| 33000/33000 [04:55<00:00, 111.55it/s]

```

```

In [0]: Xtrain.drop(['preprocessed_q1', 'preprocessed_q2'], axis=1, inplace=True)
Xtest.drop(['preprocessed_q1', 'preprocessed_q2'], axis=1, inplace=True)

```

```

In [0]: n = np.asarray(Xtrain.values)
n1 = np.asarray(Xtest.values)

```

```

In [0]: # stacking train and test data
train = hs((a,b,n)) # data with tfidf vector
test = hs((c,d,n1))

train_2 = hstack((np.asarray(a1), np.asarray(a2), n)) # data with tfidf-w2vec
vector
test_2 = hstack((np.asarray(b1), np.asarray(b2), n1))

```

```

In [0]: def predict_with_best_t(proba, threshold):
    predictions = []
    for i in proba:
        predictions.append(1 if i>=threshold else 0)
    return predictions

```

```

In [0]: def precision(arr):
    arr = arr.astype(float)
    for i in range(2):
        x = arr[0][i]+arr[1][i]
        arr[0][i] = arr[0][i]/x
        arr[1][i] = arr[1][i]/x
    return arr

def recall(arr):
    arr = arr.astype(float)
    for i in range(2):
        x = arr[i][0]+arr[i][1]
        arr[i][0] = float(arr[i][0]/x)
        arr[i][1] = float(arr[i][1]/x)
    return arr

def confussion_matrix(pred, actual, t ):

    # https://datatofish.com/confusion-matrix-python/
    d = {'y_Predicted': list(pred),
        'y_Actual': actual
        }

    df = pd.DataFrame(d, columns=['y_Actual', 'y_Predicted'])
    a = pd.crosstab(df['y_Actual'], df['y_Predicted'], rownames=['Actual'], col
names=['Predicted'])

    p = precision(np.array(a))

```



```

p = pd.DataFrame(p, columns=['0', '1'])

r = recall(np.array(a))
r = pd.DataFrame(r, columns=['0', '1'])

fig = plt.figure(figsize=(15,4))

plt.subplot(1, 3, 1)
ax = sns.heatmap(a,cbar=False, fmt='g', annot=True)
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
plt.title('confusion matrix '+t)

plt.subplot(1, 3, 2)
df_corr = p.corr()
ax = sns.heatmap(p,cbar=False, fmt='g', annot=True)
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
plt.title('precision matrix '+t)
plt.ylabel('Actual')
plt.xlabel('Predicted')

plt.subplot(1, 3, 3)
ax = sns.heatmap(r,cbar=False, fmt='g', annot=True)
bottom, top = ax.get_ylim()
ax.set_ylim(bottom + 0.5, top - 0.5)
plt.title('recall matrix '+t)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.show()
plt.close()

```

Random Model


```

In [0]: # we need to generate 9 numbers and the sum of numbers should be 1
# one solution is to generate 9 numbers and divide each of the numbers by their
sum
# ref: https://stackoverflow.com/a/18662466/4084039
# we create a output array that has exactly same size as the CV data
t = len(Ytest)
predicted_y = np.zeros((t,2))
for i in range(t):
    rand_probs = np.random.rand(1,2)
    predicted_y[i] = ((rand_probs/sum(sum(rand_probs))))[0])
print("Log loss on Test Data using Random Model",log_loss(Ytest, predicted_y, e
ps=1e-15))

predicted_y =np.argmax(predicted_y, axis=1)
confussion_matrix(Ytest, predicted_y, 'random')

```

Log loss on Test Data using Random Model 0.886734072621072



Summary

We can see that if we use the random model for our data we get the log loss of 0.887 which is very poor

Logistic Regression

```
In [0]: %%time
parameter = [10**x for x in range(-2,5)]
log_error_test = []
log_error_train = []

for i in (parameter):
    LR = SGDClassifier(alpha=i, penalty='l2', loss='log', random_state=42) # using SGDClassifier for logistic regression with loss= log
    sig_clf = CalibratedClassifierCV(LR, method="sigmoid")
    sig_clf.fit(train, Ytrain)
    predicted_train = sig_clf.predict_proba(train)
    predicted_test = sig_clf.predict_proba(test)
    log_error_test.append(log_loss(Ytest, predicted_test, labels=sig_clf.classes_, eps=1e-15))
    log_error_train.append(log_loss(Ytrain, predicted_train, labels=sig_clf.classes_, eps=1e-15))
    print(' For values of alpha = ', i, "The log loss is:", log_loss(Ytest, predicted_test, labels=sig_clf.classes_, eps=1e-15))

plt.figure(figsize = (7, 7), facecolor = None)
plt.plot(parameter, log_error_test, c='g')
for i, txt in enumerate(np.round(log_error_test,3)):
    plt.annotate((parameter[i], np.round(txt,3)), (parameter[i], log_error_test[i]))


plt.plot(parameter, log_error_train, label='Train log loss')
plt.scatter(parameter, log_error_test, label='Test log loss')
plt.plot(parameter, log_error_train, label='Train log loss')
plt.scatter(parameter, log_error_train, label='Train log loss')


plt.title('Hyper parameter v/s log loss')
plt.xlabel('alpha hyperparameter')
plt.ylabel('log loss')
plt.title('')
plt.grid()
plt.legend()

plt.show()

plt.close()
```

```
For values of alpha = 0.01 The log loss is: 0.659962416173107
For values of alpha = 0.1 The log loss is: 0.659962416173107
For values of alpha = 1 The log loss is: 0.6172207779075802
For values of alpha = 10 The log loss is: 0.6279525438922822
For values of alpha = 100 The log loss is: 0.6242902603811931
For values of alpha = 1000 The log loss is: 0.5820093471693567
For values of alpha = 10000 The log loss is: 0.6027720635127032
```





CPU times: user 1min 16s, sys: 4.26 s, total: 1min 20s

Wall time: 1min 15s


```
In [0]: best_alpha = np.argmin(log_error_test) # choosing alpha with minimum log loss
LR = SGDClassifier(alpha=parameter[best_alpha], penalty='l2', loss='log', random_state=42) # using SGDClassifier for logistic regression with loss= log
sig_clf = CalibratedClassifierCV(LR, method="sigmoid")
sig_clf.fit(train, Ytrain)
predicted_train = sig_clf.predict_proba(train)
predicted_train = [i[1] for i in predicted_train]
predicted_test = sig_clf.predict_proba(test)
predicted_test = [i[1] for i in predicted_test]

print(' For alpha = ', parameter[best_alpha], "The train loss is:", log_loss(Ytrain, predicted_train, labels=sig_clf.classes_, eps=1e-15))
print(' For alpha = ', parameter[best_alpha], "The test loss is:", log_loss(Ytest, predicted_test, labels=sig_clf.classes_, eps=1e-15))
```

For alpha = 1000 The train loss is: 0.5816603136475661

For alpha = 1000 The test loss is: 0.5820093471693568

```
In [0]: pred_train = predict_with_best_t(predicted_train, log_error_train[best_alpha])
pred_test = predict_with_best_t(predicted_test, log_error_test[best_alpha])
confusion_matrix(pred_train, Ytrain.tolist(), 'train')
confusion_matrix(pred_test, Ytest.tolist(), 'test')
```



Summary

The log loss using logistic regression is 0.5816 which is better than random model but is greater than .5 which means model performance is not good with logistic regression

Linear Svm

```
In [0]: %%time
parameter = [10**x for x in range(-2,3)]
log_error_test = []
log_error_train = []

for i in (parameter):
    LR = SGDClassifier(alpha=i, penalty='l2', loss='hinge', random_state=42) # using SGDClassifier for logistic regression with loss= log
    sig_clf = CalibratedClassifierCV(LR, method="sigmoid")
    sig_clf.fit(train, Ytrain)
    predicted_train = sig_clf.predict_proba(train)
    predicted_test = sig_clf.predict_proba(test)
    log_error_test.append(log_loss(Ytest, predicted_test, labels=sig_clf.classes_, eps=1e-15))
    log_error_train.append(log_loss(Ytrain, predicted_train, labels=sig_clf.classes_, eps=1e-15))
    print(' For values of alpha = ', i, "The log loss is:", log_loss(Ytest, predicted_test, labels=sig_clf.classes_, eps=1e-15))

plt.figure(figsize = (7, 7), facecolor = None)
plt.plot(parameter, log_error_test, c='g')
for i, txt in enumerate(np.round(log_error_test, 3)):
    plt.annotate((parameter[i], np.round(txt, 3)), (parameter[i], log_error_test[i]))


plt.plot(parameter, log_error_test, label='Test log loss')
plt.scatter(parameter, log_error_test, label='Test log loss')
plt.plot(parameter, log_error_train, label='Train log loss')
plt.scatter(parameter, log_error_train, label='Train log loss')

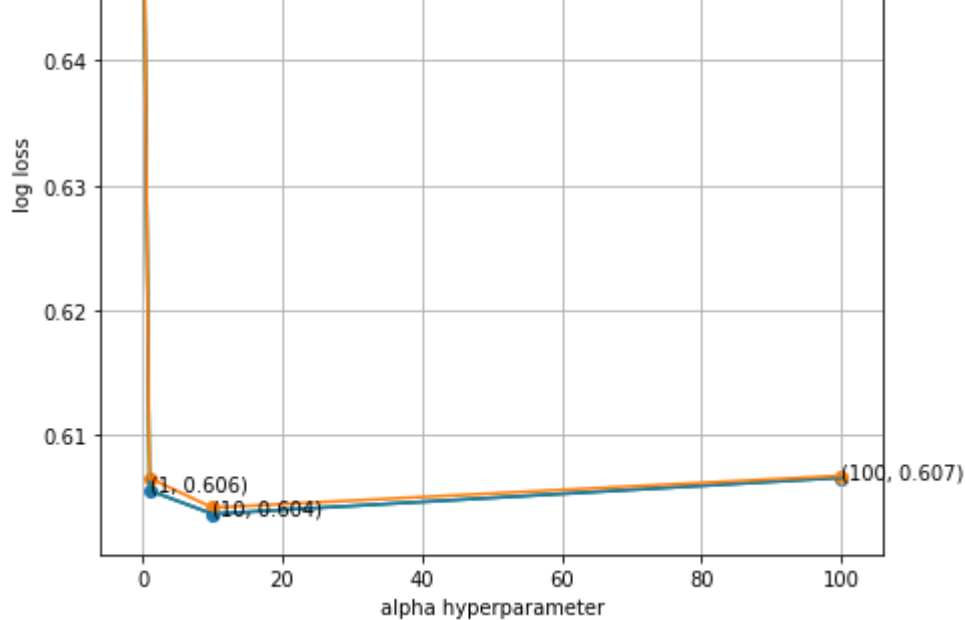
plt.title('Hyper parameter v/s log loss')
plt.xlabel('alpha hyperparameter')
plt.ylabel('log loss')
plt.title('error v/s alpha')
plt.grid()

plt.show()

plt.close()
```

```
For values of alpha = 0.01 The log loss is: 0.6573141219912
For values of alpha = 0.1 The log loss is: 0.6573141219912
For values of alpha = 1 The log loss is: 0.6055505220105439
For values of alpha = 10 The log loss is: 0.603665866397361
For values of alpha = 100 The log loss is: 0.6065704367746616
```





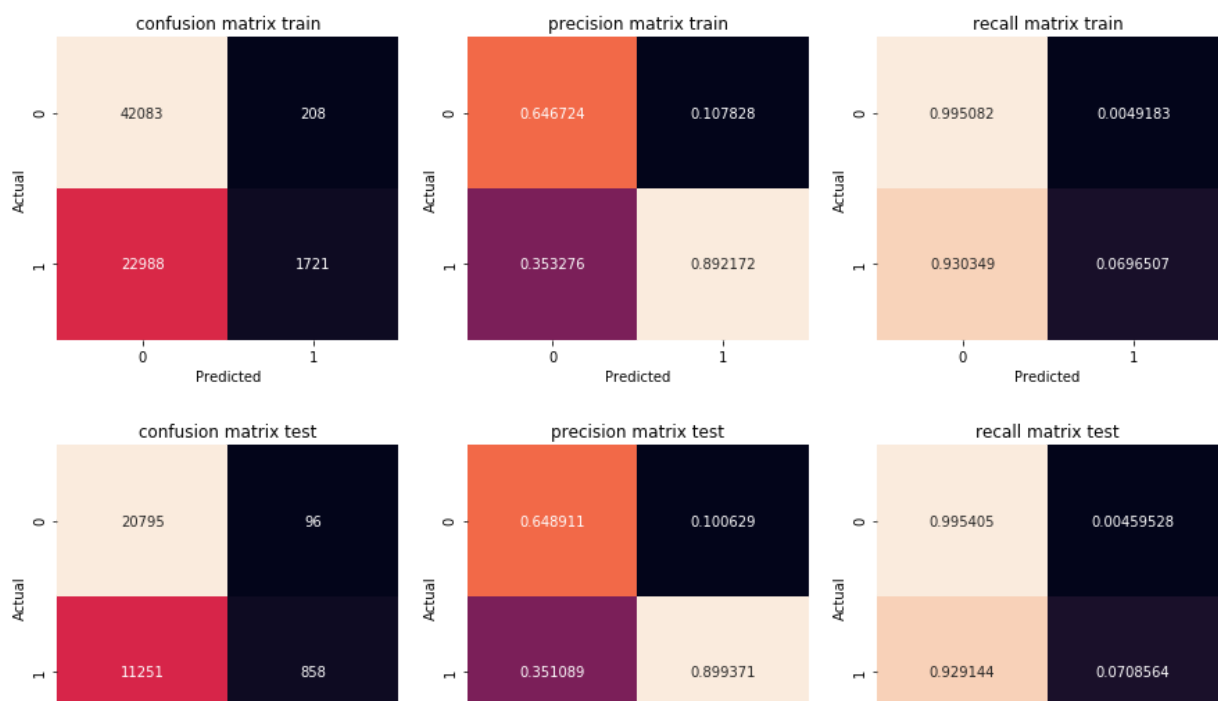
CPU times: user 42.9 s, sys: 2.99 s, total: 45.8 s
Wall time: 42.4 s

```
In [0]: best_alpha = np.argmin(log_error_test)
LR = SGDClassifier(alpha=parameter[best_alpha], penalty='l2', loss='hinge', ran
dom_state=42) # using SGDClassifier for logistic regression with loss= log
sig_clf = CalibratedClassifierCV(LR, method="sigmoid")
sig_clf.fit(train, Ytrain)
predicted_train = sig_clf.predict_proba(train)
predicted_train = [i[1] for i in predicted_train]
predicted_test = sig_clf.predict_proba(test)
predicted_test = [i[1] for i in predicted_test]

print(' For alpha = ', best_alpha, "The train loss is:", log_loss(Ytrain, predic
ted_train, labels=sig_clf.classes_, eps=1e-15))
print(' For alpha = ', best_alpha, "The test loss is:", log_loss(Ytest, predi
cted_test, labels=sig_clf.classes_, eps=1e-15))
```

For alpha = 3 The train loss is: 0.6041668550223832
For alpha = 3 The test loss is: 0.603665866397361

```
In [0]: pred_train = predict_with_best_t(predicted_train, log_error_train[best_alpha])
pred_test = predict_with_best_t(predicted_test, log_error_test[best_alpha])
confusion_matrix(pred_train, Ytrain.tolist(), 'train' )
confusion_matrix(pred_test, Ytest.tolist(), 'test' )
```



Summary

The log loss using SGDClassifier with hinge loss is 0.6041 which is poor than logistic regression model, which means model performance is not good with Linea Svm

GBDT using XGboost

```
In [0]: def plot_heatmaps(parameter, param_name, train_error, cv_error):
    # creating temporary data frame to plot heatmaps
    tmp11 = pd.DataFrame()
    tmp11[param_name[0]] = parameter[0] # k11
    tmp11[param_name[1]] = parameter[1] #K21
    tmp11['train'] = train_error # train_error
    tmp11['cv'] = cv_error #cv_error

    #https://github.com/mGalarnyk/Python_Tutorials/blob/master/Request/Heat%20M
    #aps%20using%20Matplotlib%20and%20Seaborn.ipynb
    fig, axes = plt.subplots(nrows = 1, ncols = 2, figsize = (20,20));
    pivot_table = tmp11.pivot(param_name[0], param_name[1], 'train')
    pivot_table2=tmp11.pivot(param_name[0], param_name[1], 'cv')
    axes[0].set_title('Log-loss for Train data',size=15)
    axes[0].set_xlabel(param_name[0],fontsize=15)
    axes[0].set_ylabel(param_name[1],fontsize=15)
    axes[1].set_xlabel(param_name[0],fontsize=15)
    axes[1].set_ylabel(param_name[1],fontsize=15)
    axes[1].set_title('Log-loss for cv data',size=15)
    ax = sns.heatmap(pivot_table, annot=True,annot_kws={"size": 15}, fmt='.3f',
    linewidths=.5, square = True, cmap = 'cool', cbar=False,ax=axes[0])
    bottom, top = ax.get_ylim()
    ax.set_ylim(bottom + 0.5, top - 0.5)
    ax = sns.heatmap(pivot_table2, annot=True,annot_kws={"size": 15},fmt='.3f',
    linewidths=.5, square = True, cmap = 'cool',cbar=False,ax=axes[1])
    bottom, top = ax.get_ylim()
    ax.set_ylim(bottom + 0.5, top - 0.5)
```

```
In [0]: # to train the XGBClassifier model below link has been referred
# https://www.analyticsvidhya.com/blog/2016/03/complete-guide-parameter-tuning-
# xgboost-with-codes-python/
# choosing 5 parameters to train the model

max_depth = [x for x in range(1,15, 3)]
n_estimators = [x for x in range(10,100, 10)]
learning_rate = list(np.linspace(0.01, 1, 5))
min_child_weight = list(np.linspace(0, 3, 5))
gamma = list(np.linspace(0,1,5))
```

```
In [207]: %%time
# tuning max_depth and n_estimators
parameter1 = {'max_depth': max_depth, 'n_estimators':n_estimators}

clf = XGBClassifier(booster = 'gbtree', nthread=4) # using XGBClassifier from x
gboost for
model = RandomizedSearchCV(clf, parameter1, n_iter=10, n_jobs=1,cv=5,scoring='n
eg_log_loss', return_train_score=True)
model.fit(train_2, Ytrain)
```


CPU times: user 2h 50min 21s, sys: 35.5 s, total: 2h 50min 56s

Wall time: 1h 27min 37s

```
In [209]: results1 = pd.DataFrame.from_dict(model.cv_results_)
# getting important metrics from results
train_error= results1['mean_train_score']
cv_error = results1['mean_test_score']
parameter = [results1['param_max_depth'], results1['param_n_estimators']]
param_name = ['max_depth', 'n_estimators']
print("number of data points = {}".format(len(Ytest)))

plot_heatmaps( parameter, param_name, train_error, cv_error)
```

number of data points = 33000



```
In [210]: %%time
# based on the above heat maps choosing max_depth=13 and
# tuning the min_child_weight and gamma hyper parameters
parameter2 = {'min_child_weight':min_child_weight, 'gamma':gamma}

clf = XGBClassifier(max_depth=13,n_estimators =40 ,booster = 'gbtree', nthread=
4) # using XGBClassifier from xgboost for
model = RandomizedSearchCV(clf, parameter2, n_iter=10, n_jobs=4,cv=5,scoring='n
eg_log_loss', return_train_score=True)
model.fit(train, Ytrain)


results2 = pd.DataFrame.from_dict(model.cv_results_)
# getting important metrics from results
train_error= results2['mean_train_score']
cv_error = results2['mean_test_score']
parameter = [results2['param_min_child_weight'], results2['param_gamma']]
param_name = ['min_child_weight', 'gamma']
print("number of data points = {}".format(len(Ytest)))

plot_heatmaps( parameter, param_name, train_error, cv_error)
```

number of data points = 33000

CPU times: user 1min 30s, sys: 1.41 s, total: 1min 32s

Wall time: 26min 37s





```
In [213]: %%time
# finding the suitable learning rate for the model
log_error_test=[]
log_error_train=[]

for i in learning_rate:
    clf = XGBClassifier(max_depth=10,min_child_weight=1.5,gamma=0.5,n_estimators=40,booster = 'gbtree',learning_rate= i,nthread=4) # using XGBClassifier from xgboost for
    #model = RandomizedSearchCV(clf, parameter3, n_iter=10, n_jobs=4,cv=5,scoring='neg_log_loss', return_train_score=True)
    clf.fit(train, Ytrain)
    pred_train = clf.predict_proba(train)
    pred_test = clf.predict_proba(test)
    log_error_test.append(log_loss(Ytest, pred_test, eps=1e-15, labels=clf.classes_))
    log_error_train.append(log_loss(Ytrain, pred_train, eps=1e-15, labels=clf.classes_))


#results = pd.DataFrame.from_dict(model.cv_results_)
# getting important metrics from results
train_error= log_error_train
test_error = log_error_test

plt.plot(learning_rate, train_error, label="train error")
plt.plot(learning_rate, test_error, label='test error')

plt.scatter(learning_rate, train_error)
plt.scatter(learning_rate, test_error)

plt.xlabel('learning_rate')
plt.ylabel('errors')
plt.grid()

plt.title(' error v/s learning rate')
plt.show()
plt.close()
```



CPU times: user 4min 6s, sys: 1.23 s, total: 4min 8s
Wall time: 2min 17s

In [238]: %%time

```
# after the important hyperparameters, training the final model
clf = XGBClassifier(max_depth=8,
                    min_child_weight=1.5,
                    gamma=0.5,
                    n_estimators=45,
                    booster = 'gbtree',
                    nthread=4,
                    colsample_bytree=1,
                    colsample_bylevel=1,
                    learning_rate=0.25,
                    verbosity=1) # using XGBClassifier from xgboost for GBDT
clf.fit(train_2, Ytrain, verbose=1)
log_error_test=[]

predicted_train = clf.predict_proba(train_2)
predicted_train = [i[1] for i in predicted_train ]
predicted_test = clf.predict_proba(test_2)
predicted_test = [i[1] for i in predicted_test ]
log_error_test.append(log_loss(Ytest, predicted_test, labels=clf.classes_, eps=
1e-15))
print( "The log loss train is:",log_loss(Ytrain, predicted_train, labels=clf.cl
asses_, eps=1e-15))
print( "The log loss test is:",log_loss(Ytest, predicted_test, labels=clf.class
es_, eps=1e-15))
```

The log loss train is: 0.2353068252385698

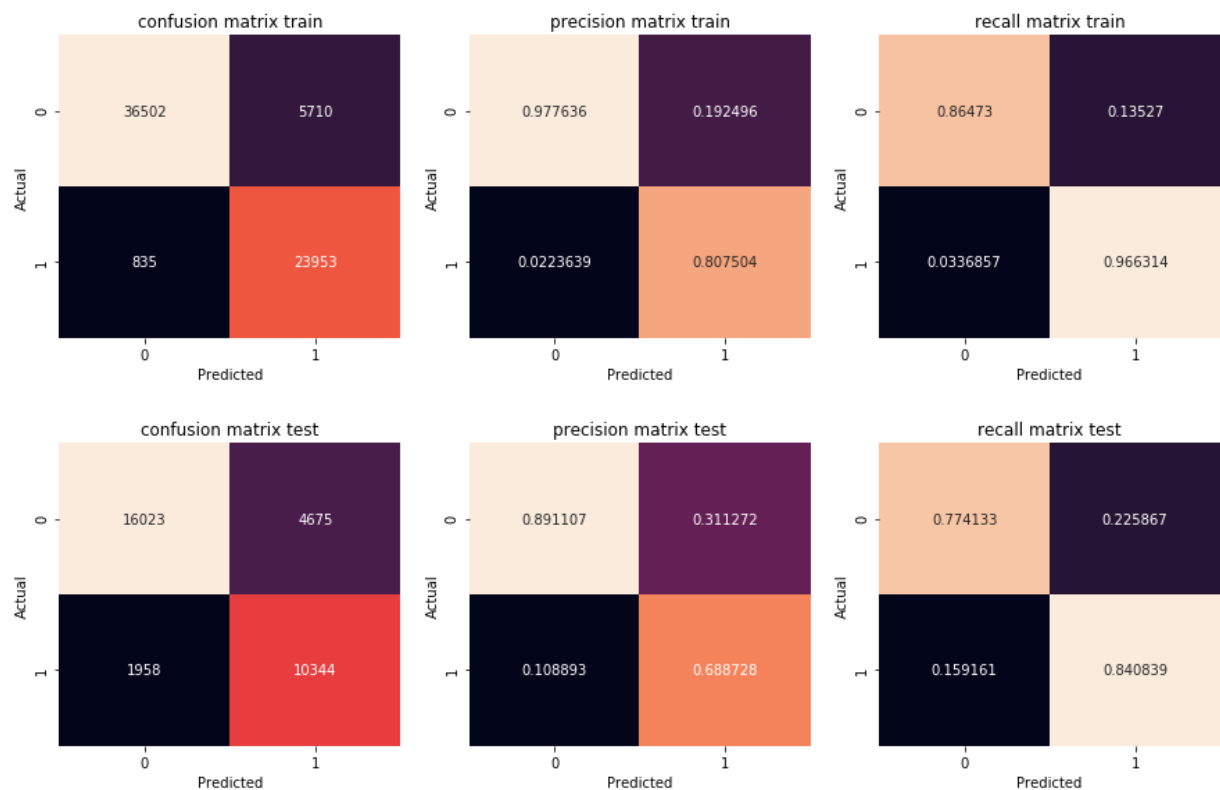
The log loss test is: 0.36935211935219076

CPU times: user 3min 37s, sys: 509 ms, total: 3min 38s

Wall time: 1min 51s

In [239]: %%time

```
log_error= log_error_test[0] # giving the value of log loss
pred_train = predict_with_best_t(predicted_train, log_error)
pred_test = predict_with_best_t(predicted_test, log_error)
confussion_matrix(pred_train,Ytrain.tolist(), 'train' )
confussion_matrix(pred_test,Ytest.tolist(), 'test' )
```



CPU times: user 1.04 s, sys: 204 ms, total: 1.24 s

Wall time: 1.02 s

Summary

We can see that log loss is 0.3693 which is much better than logistic regression model and linear svm model, also log loss is lesser than 0.5 also the precision and recall are holding well for this model

```
In [0]: from prettytable import PrettyTable
```

```
In [0]: #http://zetcode.com/python/prettytable
x = PrettyTable()
x.field_names = ["model", "Text Vetcor", "Train Loss", "Test Loss"]
x.add_row(["Logistic Regression", "TFIDF", 0.5816, 0.5820])
x.add_row(["Linear Svm", "TFIDF", 0.6041, 0.6036])
x.add_row(["GBDT", "TFIDF-W2Vec", 0.2353, 0.3693])
```

Conclusion:

```
In [246]: print(x)
```

model	Text Vetcor	Train Loss	Test Loss
Logistic Regression	TFIDF	0.5816	0.582
Linear Svm	TFIDF	0.6041	0.6036
GBDT	TFIDF-W2Vec	0.2353	0.3693

Based on the above table we can conclude that GBDT model with weighted w2vec performs the best among logistic regression and linear svm. although the the difference b/w train loss and test loss is considerable it is giving good precisoin and recall.